

JAVA Assignment: Rubik's Cube

Stefano Sandonà

Vrije Universiteit Amsterdam, Holland

1 The Rubik's Cube

A Rubik's Cube is a puzzle designed by Mr Rubik. In a Rubik's Cube, each of the faces is covered by colored stickers and in the initial setting, all the stickers of a face are of the same color and different colors are used for different faces. After performing some random twists, the goal of the game is to twist the cube until every side returns to be of a single color.

2 The Sequential Algorithm

The idea behind the sequential algorithm is simple. From a starting cube, we can twist it in each axis, row and direction (Figure 1) and from each cube of size S , after a twist, there are $6^*(S-1)$ new possible cubes. The evolution of the search can be represented as a tree. The used state space search strategy, is the **deepening depth-first search (IDDFS)**. A depth-limited search is run repeatedly, increasing the depth limit with each iteration until it reaches d , the depth of the shallowest goal state. When the bound is increased, the solution search starts again from the begin (the root), and continues until the bound level is reached. On each iteration, the algorithm visits the nodes in the search tree in the same order as depth-first search, but the cumulative order in which nodes are first visited is effectively breadth-first. The first 2 iterations of the algorithm, with the order of the cube's evaluation is shown in Figure 2. At the tree level i there are $(6^*(S-1))^i$ nodes (cubes), so there is a rapid growth of the search tree. The time complexity of IDDFS is $O(b^d)$ and its space complexity is $O(d)$, where b is the branching factor $(6^*(S-1))$ and d is the depth of the shallowest goal.




Figure 1: Possible directions and axis for twists in a 3d cube




Figure 2: The IDDFS Algorithm

3 The Parallel Algorithm

An efficient parallelization of this algorithm is not trivial because of the work distribution. At each tree level, there are $(6 * (cubeSize - 1))^i$ cubes, and this number is most likely not a multiple of the available nodes involved in the computation. Another aspect, consist on the organization of the entire work, due to the fact that after a tree bound level is reached, the termination or the progress of the work has to be checked. The next paragraph was introduced to clarify terms and concepts that will be used in the rest of the report.

3.1 The Ibis Portability Layer (IPL)

For this project was used an efficient platform for distributed computing: Ibis. This is a Java-centric communication system designed for grid computing, developed at VU University, Amsterdam. A fundamental concept in Ibis is the Pool, that is a set of one or more Ibis instances, often distributed in different machines. Each Ibis instance, that joined a specific pool, has an unique id and can cooperate with the others, running a single distributed application. With the election mechanism of Ibis, one machine can be elected as "master" and so be the "coordinator" or the central point of the entire work. Thanks to the registry mechanism, each Ibis instance can also know all the other Ibises that joined the same Pool and so communicate with them. The core of the Ibis system is the Ibis Portability layer (IPL), that is a communication library specifically designed for usage in a grid environment. It is based on setting up connections. The programmer can create send/receive ports (with a name) and connect them in a flexible way: one-to-one, one-to-many, many-to-one. In addition, the programmer can define properties of connections and ports like reliability, FIFO ordering, Proceeding in this way, after apply a certain configuration to a port, the only thing to do is to send a message through this port, the library will take care of the rest.

3.2 Cubes distribution

The work distribution, or better the cubes distribution, is the central part of the parallel implementation. The more balanced is the work among the involved nodes and the more will be the achieved performance. The major problem with a static partitioning, is that at each tree level, a number multiple of 6 of cubes are generated, and this is often not perfect divisible by the number of involved machines. A naive approach could generate the first levels of the tree and when there are enough nodes (more than the number of Ibis instances involved), distribute them as fairly as possible. However, this could lead to a big load imbalance and, as consequence, to have not good performance. For example, with a 3d cube and 5 machines, the tree will be unroll until level 1. The 12 resulting cubes would be divided into 3 3 2 2 2 and assigned to the 5 machines. This means that the first 2 machines would work 1.5 times the others. To outline the gravity of this, if m (the tree level in which lives the solution) is big enough, it could be possible that the last 3 machines worked for an entire day (24 hours), and the first two machines worked for one day and a half (36 hours). The total computational time would be 36 hours and for 12 hours the last 3 machines would be idle. The basic idea of amplify the tree until there are enough nodes is correct and logical, but the load imbalance factor must be considered.

3.2.1 Cube Weight

To avoid the load imbalance, a way for measure the total amount of work assigned to a node is needed. To do this, to each cube could be assigned a "weight", based on the amount of the next cubes that will be generated from it. This weight is represented in Equation1, where m is the tree level in which there will be the solution(s) (unknown at the begin), n is $6^{*(cubeSize-1)}$ (the number of possible children per cube) and z is the actual tree level of the cube.

$$Weight = \sum_{i=0}^{m-z} n^i \quad (1)$$

As clarification, a practical example is reported in Figure 3. In this example, that is not related to the Rubik's Cube, but just to explain the general idea, $n=2$, $z1=1$, $z2=2$, $m=4$. Applying the previous formula, the weight of the first cube (at $z1$ level) is $2^0 + 2^1 + 2^2 + 2^3 = 15$, because from it 14 cubes will be generated, so 15 (counting also the root node) will be evaluated in total. The weight of the second cube (at $z2$ level), is $2^0 + 2^1 + 2^2 = 7$.




Figure 3: Rubik's Cube Solution Tree

3.2.2 Load balancing

To have a fair ripartition of work. it is correct to assign to each Ibis instance more or less the same work to do. Given the fact that the number of cubes on a tree level often is not a multiple of the nodes involved on the computation, there will be load imbalance at a certain point. Also distributing the nodes at a specific tree level as fairly as possible, there will be one or more Ibis instances with one more cube than the others, for a gap of maximum one cube (e.g. with 9 cubes and 4 machines, 3 2 2 2). Accepting this, the best thing to do, is to select the correct tree level in which have this load imbalance, in order to limit the deriving damages. Looking at Equation1, at the begin of the computation n is known because it is the size of the cube, m is unknown, is part of the solution, and z is the value that we want to find. As just said, given enough cubes (more than the number of Ibis instances), the maximum difference of assigned cubes per node is one. This is possible using a simple *for* construct (Listing 1).

```

int sum = 0;
int rem = numCubes % numProcs;
for (i = 0; i < numprocs; i++) {
    cubes_per_proc[i] = numCubes / numProcs;
    if (rem > 0) {
        cubes_per_proc[i]++;
        rem--;
    }
    displs[i] = sum;
    sum += cubes_per_proc[i];
}

```

Listing 1: fair distribution of cubes

At a certain tree level z , there are n^z cubes and the relative weight per cube is the one shown in Equation1. With the previous code (Listing 1), calling k the number of machines involved in the

computation, the lowest amount of work assigned per slave, is shown in Equation 2, while the highest is shown in Equation 3.

$$LW = \sum_{i=0}^{m-z} n^i * \lfloor \frac{n^z}{k} \rfloor \quad (2)$$

$$HW = \sum_{i=0}^{m-z} n^i * \lfloor \frac{n^z}{k} \rfloor + \sum_{i=0}^{m-z} n^i \quad (3)$$

To reduce the load imbalance as much as possible, the two amounts of work should differ as little as possible. That means that $HW*100/LW$ has to be as near to 100 as possible, for example lower than 101, that means that a maximum of 1% of load imbalance is permitted. Solving the inequality shown in Equation 4, the solution is the one shown in Equation 5.

$$\frac{(\sum_{i=0}^{m-z} n^i * \lfloor \frac{n^z}{k} \rfloor + \sum_{i=0}^{m-z} n^i) * 100}{(\sum_{i=0}^{m-z} n^i * \lfloor \frac{n^z}{k} \rfloor)} < 101 \quad (4)$$

$$\lfloor \frac{n^z}{k} \rfloor > 100 \quad (5)$$

At the begin n (the size of the cube) is known and also k (the number of Ibises that joined the pool), so the value to find is z . In addition to reduce the load imbalance, it's useful to reduce the initial tree unroll at the minimum. For this reason, the final value to find is the one shown in Equation 6.

$$\min_{z>0} \lfloor \frac{n^z}{k} \rfloor > 100 \quad (6)$$

3.2.3 Work split

Once found the correct value of z , this is the tree level in which to have an imbalance of 1 cube, isn't such a big risk for the performance. To reduce the unroll, the tree is extended (A) until the level $z-1$, if at that level there are enough nodes (more than the Ibis instances), otherwise (B) until the level z . If we are in case A, the $n^{(z-1)}$ resulting nodes are divided into equal parts to the machines ($n^{(z-1)}/k$ cubes per machine). If some nodes have left out (k is not a divisor of the number of nodes of this level), these are expanded to the next tree level (Figure 4), otherwise the splitting phase is terminated (Figure 5). If we have expanded some nodes, we try to split their children among the Ibis instances. Until they are less than the number of ibis instances we generate another level of the tree from them (Figure 6). When they are enough, we split them as fairly as possible (using the approach shown in Listing 1). If we are in case B we just split the $n^{(z)}$ nodes as fairly as possible (Listing 1). In Figure 4, Table 1 there is an example of work ripartition among 11 Ibis instances, in Figure 5, Table 2 among 16 Ibis instances and in Figure 6, Table 3 among 13 Ibis instances.




Figure 4: Cubes distribution - Ibis instances=11, Size=3, Twists=11, Bound=7

Ibis Instance	0	1	2	3	4	5	6	7	8	9	10	Total
Cubes of level 2	13	13	13	13	13	13	13	13	13	13	13	143
Cubes of level 3	2	1	1	1	1	1	1	1	1	1	1	12

Table 1: Cubes distribution - Ibis instances=11, Size=3, Twists=11, Bound=7




Figure 5: Cubes distribution - Ibis instances=16, Size=3, Twists=11, Bound=7

Ibis Instance	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	Total
Cubes of level 2	14	14	14	14	14	14	14	14	14	14	14	14	14	14	14	14	144

Table 2: Cubes distribution - Ibis instances=16, Size=3, Twists=11, Bound=7




Figure 6: Cubes distribution - Ibis instances=13, Size=3, Twists=11, Bound=7

Ibis Instance	0	1	2	3	4	5	6	7	8	9	10	11	12	Total
Cubes of level 2	11	11	11	11	11	11	11	11	11	11	11	11	11	143
Cubes of level 3	0	0	0	0	0	0	0	0	0	0	0	0	0	0
Cubes of level 4	12	11	11	11	11	11	11	11	11	11	11	11	11	144

Table 3: Cubes distribution - Ibis instances=13, Size=3, Twists=11, Bound=7

In Table 4, are reported the numbers of cubes evaluated by each Ibis instance, when the algorithm's bound is 7, with a Cube of size 3. These numbers were measured running the implementation with different number of Ibises. The Expanded Cubes, are the ones from which are generated other cubes during the initial splitting work phase while the Assigned Cubes are the ones effectively evaluated. As shown, for every launching configuration, the load imbalance is always under the 1%.

3.3 Ibis setting

Each node involved in the computation is an Ibis instance with an unique identifier. When a tree level bound is reached, each Ibis instance has to communicate its results (if a solution is founded on the

Table 4: Cubes distribution - Size=3, Twists=11, Bound=7

Ibis Instance	Assigned Cubes	Expanded Cubes	Total Cubes	Load Imbalance
2	19544622	1	39089245	0%
4	9772311	1	39089245	0%
5	7804273(2), 7826894(3)	17	39089245	0.28%
8	4886154	13	39089245	0%
11	3551510(10), 3574131(1)	14	39089245	0.63%
13	3006718(12), 3008603(1)	26	39089245	0.06%



(a) Results communication - Many to 1

(b) Work status communication - 1 to Many

Figure 7: Send (circle) and Receive (square) Ports organization

assigned subtree and, in the case, how many) and has to decide if continue with the next bound or not. The most obvious solution, is to adopt a "master slave" approach and use the "master" Ibis instance to collect the partial results and decide if the work is finished or continue with the next bound. To do this 4 ports are needed. On the master side 2 ports are needed, one to receive the results from the slaves and one to send if the the work is finished. Specularly, on the slaves side, 2 ports are needed, one to communicate the solutions found and one to receive if the if the next bound has to be evaluated or the work is finished. A graphic representation is presented in Figure 7.

3.3.1 Weaknesses of the implementation

3.4 Results